

naiveBayes

September 25, 2021

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
import seaborn as sns
warnings.filterwarnings('ignore')

from IPython.display import set_matplotlib_formats
set_matplotlib_formats('pdf', 'svg')

from pandas import DataFrame
from numpy import arange
from pandas import read_csv
from matplotlib import pyplot

from scipy.stats import multivariate_normal as mvn

data=pd.read_csv('/content/drive/MyDrive/exNB.csv', header=None)

#print(data.head())

#print(data.describe())

X = data.to_numpy()

#print(X)

y= X[:, -1]

X=X[:, :-1]

#print(X)
#print(y)

#print(X[y==1, :])
```

```

plt.figure()
plt.hist(X[y==1,0], label="Male", alpha=.5)
plt.hist(X[y==0,0], label="Male", alpha=.5)

plt.legend()

plt.figure()

plt.hist(X[y==1,1], label="Male", alpha=.5)
plt.hist(X[y==0,1], label="Male", alpha=.5)
plt.legend()

plt.figure()

plt.figure()
plt.scatter(X[:,0],X[:,1], c=y, alpha=.25)

```

[1]:

```

[2]: class GaussNB():
    def fit(self , X, y, epsilon= 1e-1):
        self.likelihoods = dict()
        self.priors = dict()
        self.K = set(y.astype(int))

        for k in self.K:

            X_k = X[y==k,:]

            self.likelihoods[k]={"mean": X_k.mean(axis=0), "cov":X_k.var(axis=0) +
→epsilon}
            self.priors[k]=len(X_k)/len(X)

    def predict(self, X):
        N, D = X.shape
        P_hat = np.zeros((N, len(self.K)))

        for k, l in self.likelihoods.items():
            P_hat[:,k] = mvn.logpdf(X, l["mean"], l["cov"]) + np.log(self.priors[k])

        return P_hat.argmax(axis=1)

```

[3]: gnb = GaussNB()

```
gnb.fit(X,y)
```

```
y_hat = gnb.predict(X)
```

```

plt.figure()
plt.scatter(X[:,0],X[:,1], c=y_hat, alpha=.25)

[4]: def accuracy(y,y_hat):
      return np.mean(y==y_hat)

[5]: accuracy(y,y_hat)

[5]: 0.988

[5]: 

[6]: X_new=np.asarray([[68,150]])

[7]: yh_new=gnb.predict(X_new)

[8]: accuracy(y,yh_new)

[8]: 0.5

[9]: data1= pd.read_csv('/content/drive/MyDrive/xor.csv')
      data2= pd.read_csv('/content/drive/MyDrive/MNIST_train.csv')
      data2.shape

[9]: (60000, 787)

[10]: #data2=data2.fillna(data2.mean())

[11]: #####start messing with MNIST#####
      data2=data2.iloc[:, 2:]

      #####Reciprcal Transform works if inf taken_
      →away#####

      #data2=1/(data2)
      #data2['labels']=1/(data2['labels'])

      data2=data2/255
      data2['labels']=data2['labels']*255

      #####Square root doesnt work at all:
      →ERROR#####

      #data2=pow(data2,1/2)
      #data2['labels']=pow(data2['labels'],2)
      #data2['labels']=data2['labels'].astype('category')

      #data2=data2.replace(np.inf, 0)
      #data2=data2.replace(np.nan, 0)

      print(data2.shape)
      data2.head(10)

```

(60000, 785)

```
[11]: labels  0    1    2    3    4    5  ...  777  778  779  780  781  782  783
      0    5.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      1    0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      2    4.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      3    1.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      4    9.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      5    2.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      6    1.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      7    3.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      8    1.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
      9    4.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

[10 rows x 785 columns]

```
[12]: x2=data2.to_numpy()
```

```
[12]:
```

```
[13]: y2 = x2[:, 0]
      print(y2)
      x2= x2[:, 1:]

      print(x2.shape)
```

[5. 0. 4. ... 5. 6. 8.]

(60000, 784)

```
[14]: #plt.figure()
      #plt.scatter(x2[:,0],x2[:,1], c=y2, alpha=.25)
```

```
[15]: xgnb = GaussNB()

      xgnb.fit(x2,y2)

      y_hat1 = xgnb.predict(x2)

      #plt.figure()
      #plt.scatter(x2[:,0],x2[:,1], c=y_hat1, alpha=.25)
```

```
[16]: accuracy(y2,y_hat1)
```

```
[16]: 0.7651333333333333
```

```
[17]: class GaussBayes():
      def fit(self, X, y, epsilon=1e-1):
          self.likelihoods=dict()
          self.priors= dict()

          self.K = set(y.astype(int))
```

```

    for k in self.K:
        X_k = X[y==k,:]
        N_k, D = X_k.shape
        mu_k=X_k.mean(axis=0)
        self.likelihoods[k] = {"mean":X_k.mean(axis=0), "cov":(1/(N_k-1))*np.
→matmul((X_k-mu_k).T, X_k - mu_k) +epsilon*np.identity(D)}

        self.priors[k]= len(X_k)/len(X)

    def predict(self, X):
        N,D=X.shape
        P_hat = np.zeros((N,len(self.K)))

        for k,l in self.likelihoods.items():
            P_hat[:,k] = mvn.logpdf(X, l["mean"],l["cov"]) + np.log(self.priors[k])

        return P_hat.argmax(axis=1)

```

```
[18]: gbayes = GaussBayes()
```

```
[19]: gbayes.fit(x2,y2)
```

```
[20]: y_hat_GB=gbayes.predict(x2)
```

```
[21]: #plt.figure()
      #plt.scatter(x2[:,0],x2[:,1], c=y_hat_GB, alpha=.25)
```

```
[22]: accuracy(y2,y_hat_GB)
```

```
[22]: 0.9549333333333333
```

```
[23]: #####Run Test Data#####
      data_test= pd.read_csv('/content/drive/MyDrive/MNIST_test.csv')
```

```
[24]: data_test=data_test.iloc[:, 2:]
      data_test.head()
```

```
[24]: labels  0  1  2  3  4  5  6  7  ...  775  776  777  778  779  780  781  782
783
0          7  0  0  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
0
1          2  0  0  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
0
2          1  0  0  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
0
3          0  0  0  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
0
4          4  0  0  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
0
```

[5 rows x 785 columns]


```
[25]: X3=data_test.to_numpy()
```

```
[26]: def show_me(X):  
      plt.imshow(X.reshape(28,28))  
  
      def show_me_allmean(X,y,k):  
          show_me(sum(X[y==k,:]/len(X[y==k,:])))
```

```
[27]: y3 = X3[:, 0]  
      print(y3)  
      X3= X3[:, 1:]  
  
      print(X3.shape)
```

[7 2 1 ... 4 5 6]
(10000, 784)

```
[28]: show_me(X3[96])
```



```
[29]: data_test.head(10)
```

```
[29]: labels 0 1 2 3 4 5 6 7 ... 775 776 777 778 779 780 781 782  
783  
0      7 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0  
0  
1      2 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
```

```

0
2      1 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
3      0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
4      4 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
5      1 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
6      4 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
7      9 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
8      5 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0
9      9 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0
0

```

[10 rows x 785 columns]

```
[30]: data_test.tail(10)
```

```

[30]:      labels  0  1  2  3  4  5  6  ...  776  777  778  779  780  781  782  783
9990      7  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9991      8  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9992      9  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9993      0  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9994      1  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9995      2  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9996      3  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9997      4  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9998      5  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0
9999      6  0  0  0  0  0  0  ...    0    0    0    0    0    0    0    0

```

[10 rows x 785 columns]

```
[31]: y_hat_GB2=gbayes.predict(X3)
```

```
[32]: accuracy(y3,y_hat_GB2)
```


```
[32]: 0.9423
```

```

[33]: plt.figure(figsize=(10,7))
      y_actu = pd.Series(y3, name='Actual')
      y_pred = pd.Series(y_hat_GB2, name='Predicted')
      cm = pd.crosstab(y_actu, y_pred)
      ax = sns.heatmap(cm, annot=True, fmt="d")
      plt.ylabel('True label')
      plt.xlabel('Predicted label')

```

[33]: Text(0.5, 42.0, 'Predicted label')



```
[34]: isMatch= y_actu==y_pred
      isMatch
```

```
[34]: 0      True
      1      True
      2      True
      3      True
      4      True
      ...
      9995   True
      9996   True
      9997   True
      9998   True
      9999   True
      Length: 10000, dtype: bool
```

```
[35]: misLabel = pd.concat([y_actu, y_pred,isMatch], axis=1)
      misLabel.to_csv('isMatch.csv')
```



```
[36]: %%capture
      !jupyter nbconvert --to PDF "/content/drive/MyDrive/Colab Notebooks/niaveBayes.
      →ipynb"
```